

Deep Learning Techniques for Music Generation

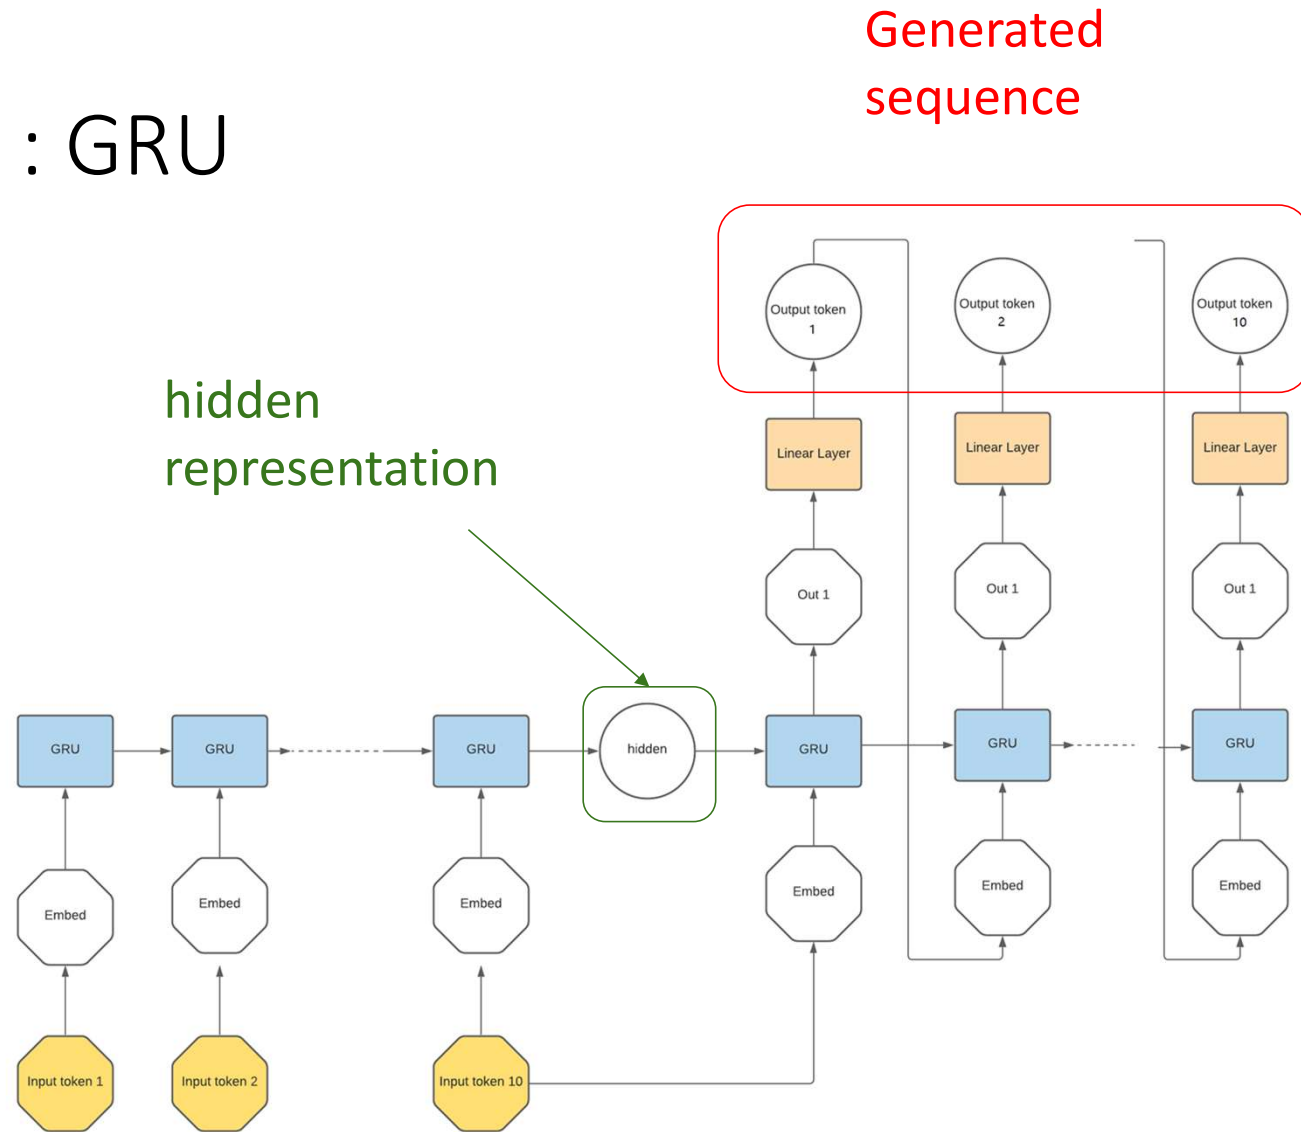
Yassine Abbahaddou

Elias Aouad

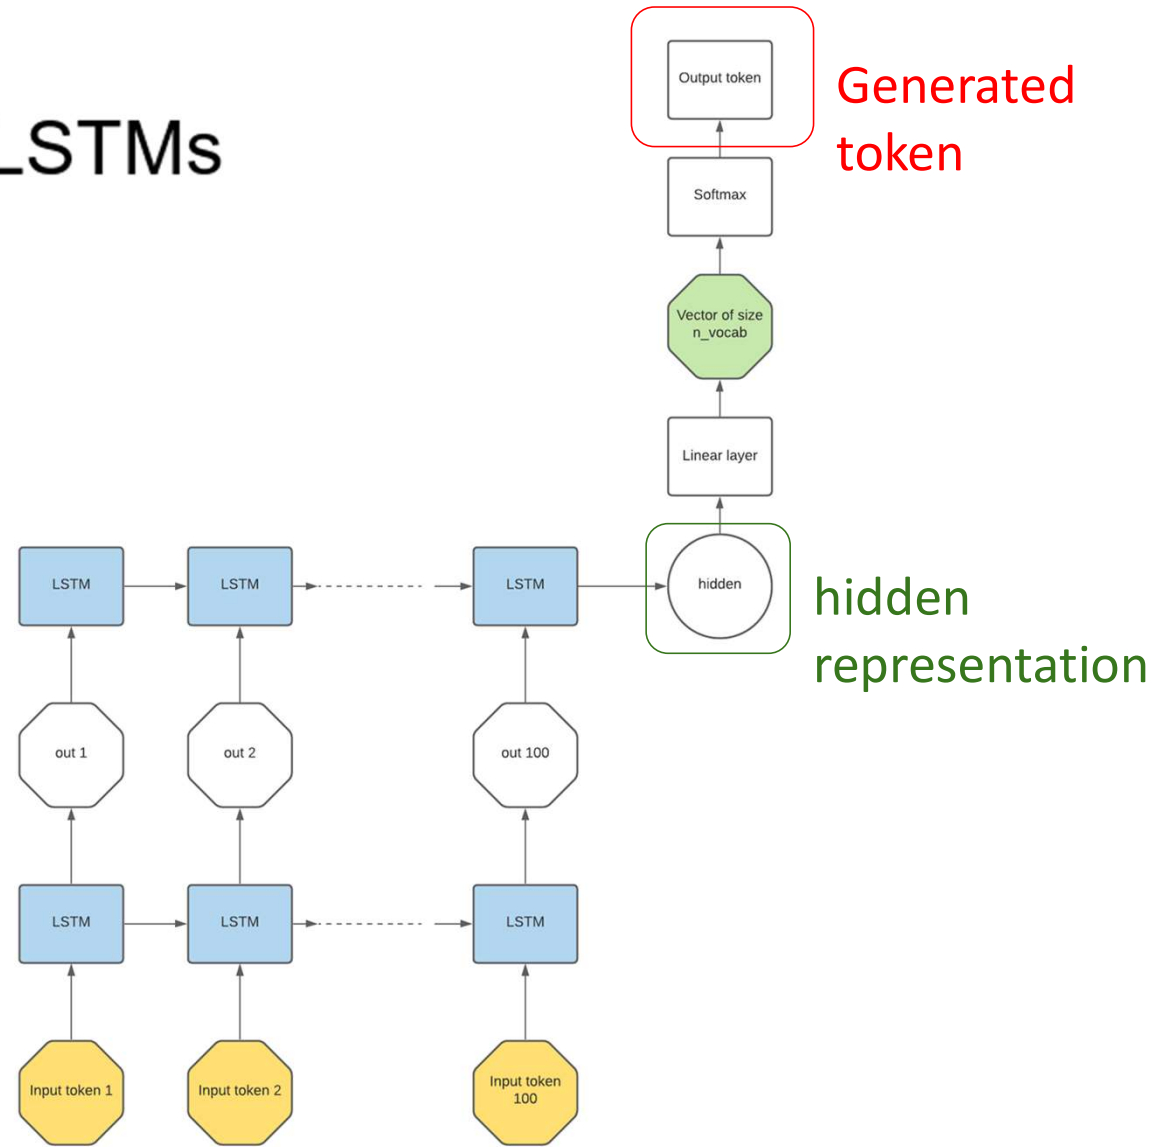
Introduction

- Goal : Generate new music using deep learning methods
- Two ways to treat this problem :
 - Use music as a raw audio
 - Encode music into textual notations
- Evaluation : Use BLEU score to rate music / Discriminator error (GAN only)

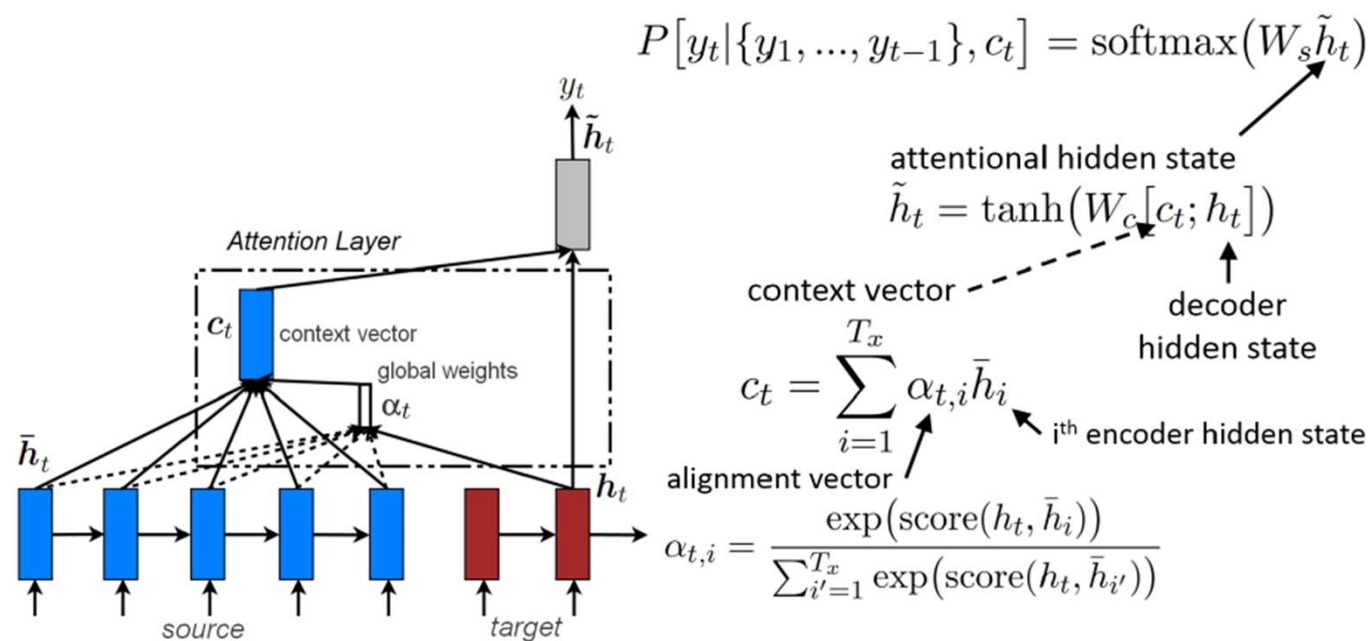
Method 1 : GRU



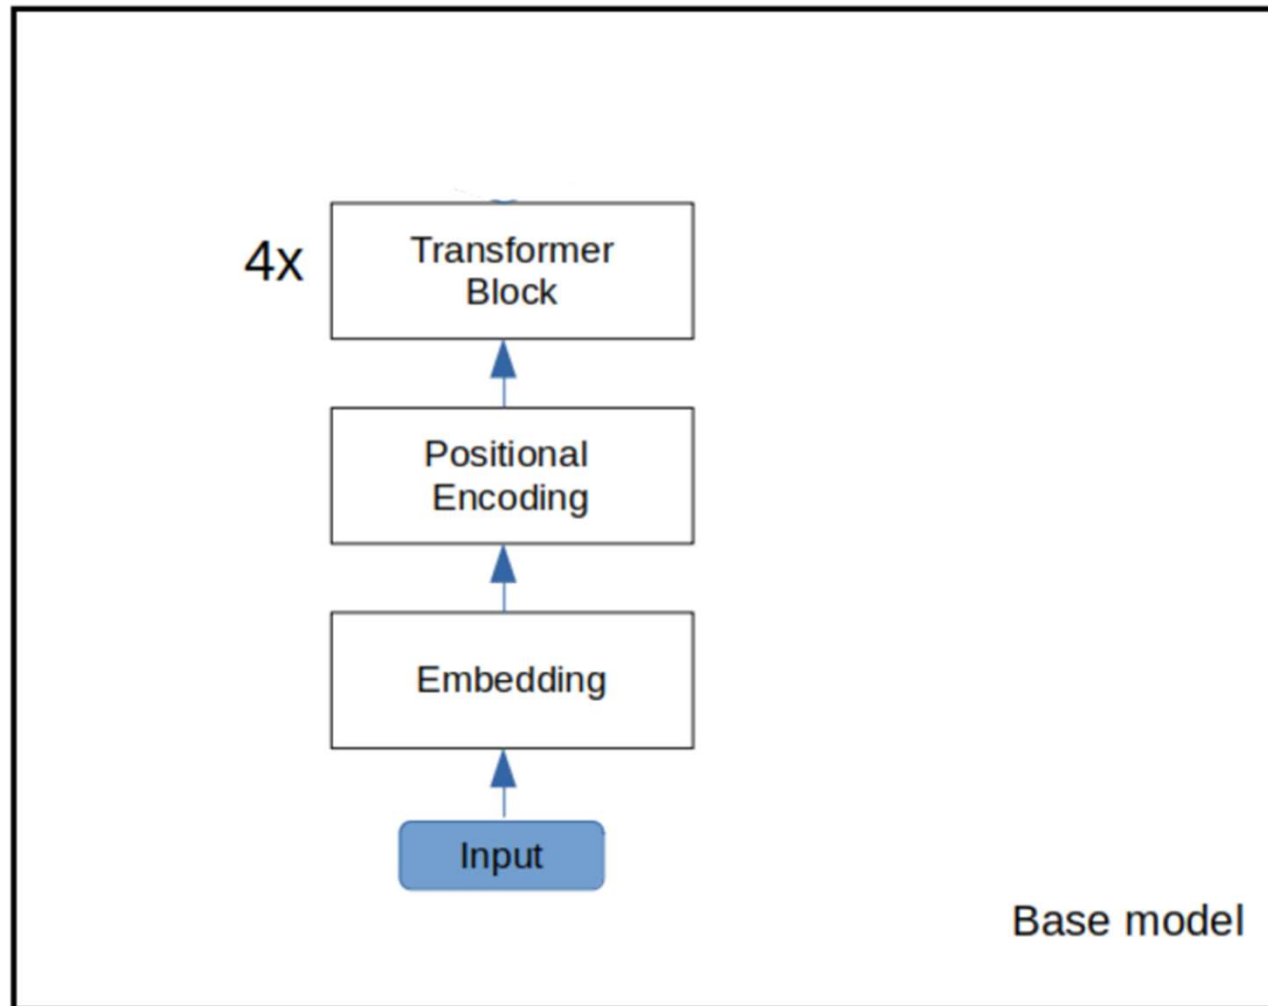
Method 2 : 2-layer LSTMs




Method 3 : GRU + attention mechanism




Method 4 : transformer Decoder



Method 5 : GAN



Method 5 :GAN



Metric : BLEU score

$$BLEU = \frac{\sum_{S \in C} \sum_{N\text{-gram} \in S} Count_{matched}(N\text{-gram})}{\sum_{S \in C} \sum_{N\text{-gram} \in S} Count(N\text{-gram})}$$

Results

Method	BLEU score	Discriminator	Human rating
GRU	80%	-	6.5 (± 3)
2-layer LSTM	71%	-	8.5 (± 1)
RNN+attention	62%	-	7.5 (± 1.7)
Transformer Decoder	57%	-	9 (± 1)
GAN	-	0%	9 (± 0.4)

Conclusion

- We mainly relied on the latest advances in the NLP field.
- The most challenging task in music generation is the evaluation.
- In terms of results, almost all the methods produced music similar to real ones.

References

- [1] R. W. Allen Huang. Deep learning for music
- [2] I. J. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio. Generative Adversarial Networks. June 2014.
- [3] M. Post. A call for clarity in reporting bleu scores. In WMT , 2018.
- [4] A. S. B. K. Rachel Manzelli, Vijay Thakkar. Conditioning deep generative raw audio models for structured automatic music.
- [5] D. Rubino. Nottingham music database.
- [6] A. van den Oord et al. Wavenet: A generative model for raw audio.
- [7] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. 2017.

Generated music with : transformer Decoder

